


Pruned subnetworks versus same-size dense models



Filled markers are means over three or more seeds and the bars are ± 1 sample standard deviation; hollow markers are a SINGLE run and carry no bar. IMP-150 and the cm=1.8 source are single runs, so their positions are one draw from a distribution whose pooled spread is 1.29 pp. Dotted line: the organisers' published baseline over THEIR five runs -- reproduction evidence, not a comparison point; the star is the source model this project actually pruned.

Shaded region: no dense model exists at a pruned size there, so the two deepest pruned points are compared against base=16, which is larger. Neither the SNIP nor the DSP mean curve is monotonic; read the shape, not a point.