


The compute axis: only structured pruning moves left



- dense, channels_multiplier sweep
- dense, base_channels sweep
- - DSP, real MACs (n=3)
- IMP-150, real MACs (constant, n=1)
- IMP-150, hypothetical -- unreachable (needs sparse kernels)
- ◇ SNIP, real MACs (constant, n=3)
- ◇ SNIP, hypothetical -- unreachable (needs sparse kernels)
- ☆ cm=1.8 source model, n=1 (both sweeps' top end)
- mean of ≥ 3 seeds, bar = ± 1 s.d.
- single seed (no bar)

Filled markers with bars are what the deployed model actually executes, averaged over three seeds; IMP's and SNIP's sit in a single vertical stack because masking changes no tensor shape, and IMP's are hollow because IMP is a single run. Hollow, dotted markers are what their surviving weights WOULD cost if every zero could be skipped -- unreachable without sparse kernels, which are out of scope here, and never to be quoted unlabelled.

The two dense families cross: base=24 has more parameters than cm=1.0 but far fewer MACs, because cm only widens the deep narrow stages where almost no compute lives.