The compute axis: only structured pruning moves left

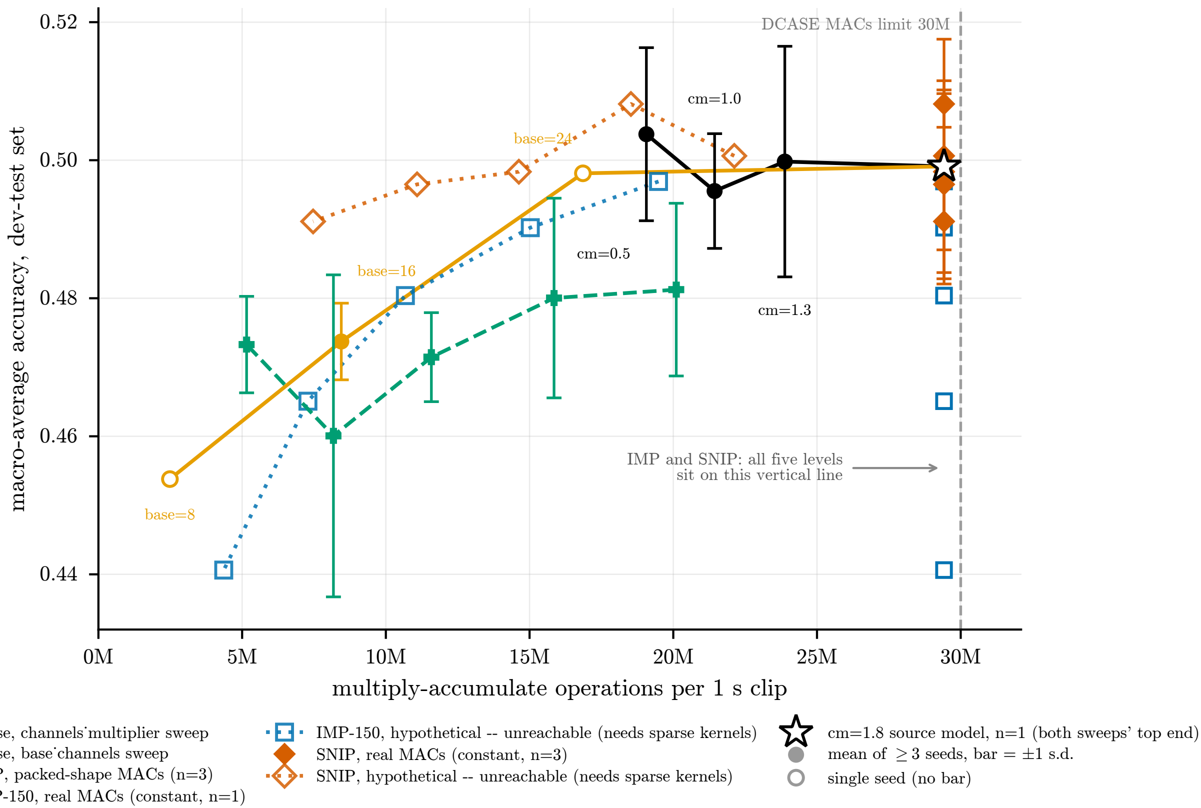

Filled markers with bars are what the tensor shapes cost: dense execution for IMP, SNIP and the dense references, calculated packed shapes for DSP, whose packing is not implemented. Averaged over three seeds; IMP's and SNIP's sit in a single vertical stack because masking changes no tensor shape, and IMP's are hollow because IMP is a single run. Hollow, dotted markers are what their surviving weights WOULD cost if every zero could be skipped -- unreachable without sparse kernels, which are out of scope here, and never to be quoted unlabelled.
The two dense families cross: base=24 has more parameters than cm=1.0 but far fewer MACs, because cm only widens the deep narrow stages where almost no compute lives.