


IMP's retraining requirement grows with sparsity



The three IMP curves are the same pruning criterion at three retraining budgets, so the vertical spread between them is compute, not method. It is negligible above 81 KB and worth 7.05 pp at 22.41 KB -- more than twice the 1.29 pp pooled seed scatter, which is why this comparison survives at one seed. All three IMP curves are hollow because each is a SINGLE run; only the dense reference here has seed replication, and its bars are ± 1 s.d. where it does. IMP-150 matches SNIP's per-level training budget exactly, which is what removes the compute confound from Figure 1.