


Every individual run behind the means in Figures 1 and 2



The small number under each level is that level's sample s.d. in percentage points. Hollow circles are the three individual seeds at each level, spread horizontally only to stop them overplotting; the thin vertical line is their min-max range and the filled marker is the mean with ± 1 sample s.d. Both panels share the accuracy axis, so the two spreads are directly comparable. DSP's are wider on average and its worst level is far wider than SNIP's worst, which is expected: a SNIP seed changes only which weights the mask keeps, while a DSP seed re-runs phase A and therefore trains a different architecture. Note what is NOT visible here: a level at which either method reliably fails. Each single-seed curve has one isolated low point and they land at different sizes.